A Distributed Real-Time Intrusion Detection System on a Jetson Orin Nano Super Edge Cluster using Kafka and PySpark

Thai Nguyen Vu[1] Dung Van Bui[2] Nhut Tri Do[2] Van Du Nguyen[3]

SOICT 2026

[1]Campus in Ho Chi Minh City, University of Transport and Communications, Ho Chi Minh City, Vietnam
[2]University of Information Technology, VNU-HCM, Ho Chi Minh City, Vietnam
[3]Faculty of Information Technology, Nong Lam University, Ho Chi Minh City, Vietnam

I. PROBLEM AND CONTRIBUTIONS

Edge intrusion detection: three open gaps

Detectors now *fit* on single-board hardware — but a monolithic one classifies **every** flow and saturates an 8 GB ARM board.

What the literature does not tell us

1. Everything is measured on a **single device**.
2. Throughput, latency and **energy** are never reported together.
3. Offline scores come from **leaky protocols**.

Our question

How should inference be *placed* across a constrained two-node edge cluster?

Our answer

Split it: a cheap **autoencoder gate** on board #1, the heavy **PySpark classifier** on board #2, joined by a Kafka topic.

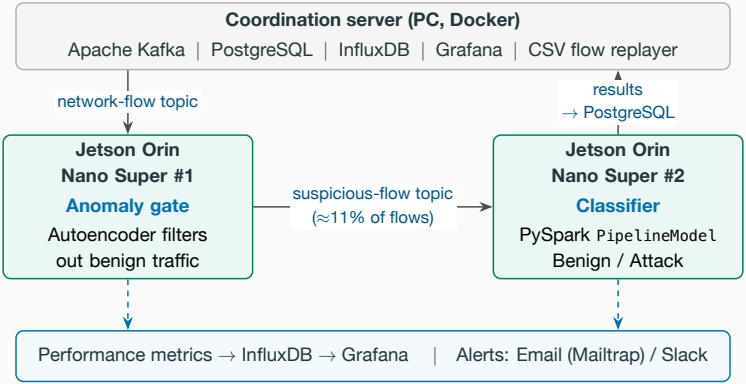
Contributions

The parts are standard (autoencoder, Kafka, Spark MLlib). **No new detector**: the contribution is what we measured on real edge hardware.

1. A **placement study** of a two-stage IDS on a two-board Jetson cluster: single node, pipeline split on one and on two boards, horizontal Kafka scaling, Spark standalone cluster.
2. An **end-to-end evaluation**: sustained capacity, latency with queueing, energy per verdict, and **precision / recall / F1 after the gate** — on CICIDS2017 and on CSE-CIC-IDS2018 replay.
3. A **same-model engine baseline**: PySpark vs. NumPy vs. ONNX Runtime on the deployed forest.
4. **Leakage-aware provenance**: the deployed model is the one selected by our offline Spark ML evaluation under a leakage-free protocol.

II. SYSTEM AND EVALUATION SETUP

Split-deployment architecture



The gate scores all 77 features; its threshold is fixed *offline* on benign validation data, never on the test set. Batches of 20, or after 1 s.

CICIDS2017: Sharafaldin et al., ICISSP 2018 · Kafka: Kreps et al., NetDB 2011 · Spark: Zaharia et al., HotCloud 2010.

Deployment modes and what we measure

Single node full pipeline on one board, gate off or on.

A: Pipeline split gate on #1, classifier on #2 (our design).

A1: Split, one board gate and classifier as two processes on #2.

B: Horizontal both boards run the full role (gate off).

C: Spark cluster host as master, both boards as workers.

Every flow is followed from send to verdict, queueing included. **Sustained capacity**: the highest rate at which every flow gets its verdict within 10 s of the load.

Setup

2× Jetson Orin Nano Super (6-core A78AE, 8 GB, ≈67 TOPS)

PySpark 3.4.1, Kafka 3.5, Wi-Fi LAN

10,000 test-split flows, 15% attacks + 10,000 CSE-CIC-IDS2018 flows

Rate sweep 1–40 flows/s, 2 runs each

Energy via `tegrastats`

III. RESULTS AND CONCLUSION

Result 1: decoupling the stages sustains 7–9× the load

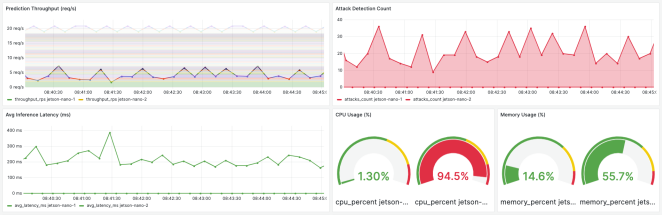
Mode	Flows/s	e2e p95 (s)	J / verdict
Single node	2.5	11.3	3.04
Single node + gate	≈2.5	7.8	3.04
B: Horizontal	3.3	8.2	3.20
C: Spark cluster	9–17	5.2	1.15
A1: Split, one board	**17.8**	5.3	**0.32**
A: Split	**22.4**	**7.1**	**0.60**

Flows/s: sustained, every flow done within 10 s. **e2e p95**: send to verdict, queueing included.

- **The gate inside the classifier's process gains nothing.**
- **As a separate process on the same board: 7×**, at 0.32 J per verdict.
- A second board adds a quarter (22.4) at twice the energy.

Decoupling, not the second board, pays: one board running the split is the efficient choice; a second board buys headroom.

Result 2: the load asymmetry is visible live



At peak load, split mode (full 10,000-flow replay, deployed model):

- gate node: **1.3% CPU**
- classifier node: **94.5% CPU**

The classifier dominates the compute budget — which is exactly why moving it off the ingest path pays.

Result 3: what the gate costs, end to end

Mode	Attacks	Missed: gate / clf	FP	F1
A: full replay	1,518	463 / 0	0	**0.820**
A1: split, one board	882	278 / 0	0	0.813
C: Spark cluster	381	115 / 0	0	0.822
Gate off	242	– / 0	0	1.000

Final verdicts of the deployed pipeline against the replay labels. Full replay: all 10,000 flows once, 20 flows/s for 500 s. Classifier alone, offline: F1 0.996.

- **Precision 100%**: no false alarm on 8,483 benign flows.
- **Recall 69.5%**: every missed attack was *skipped by the gate*; the classifier made no error on what it received.
- The gate's threshold alone sets the end-to-end F1.

The capacity gain is paid for in recall: the gate is a **tunable trade-off** (operating curve in backup), not a free filter.

Result 4: another network — CSE-CIC-IDS2018

Replayed flows	CICIDS2017	CSE-CIC-IDS2018
Gate forwards	11%	**30%**
Benign it forwards	0.6%	28%
Sustained (flows/s)	22.4	**≈8**
Recall / F1	69.5% / 0.82	**0.8% / 0.016**
Classifier alone, recall	99.8%	1.8%

Same deployed split, models trained on CICIDS2017, 10,000 flows each, replayed once. Precision 100% in both.

- The gate flags new benign traffic as anomalous: the classifier gets **3× the load**.
- The classifier itself misses almost every 2018 attack.
- Offline: 1% labelled target flows restore F1 to 0.92.

Neither detection *nor capacity* transfers: a new network needs a **recalibrated gate** and a **retrained classifier**.

Result 5: the price of one unified stack

We serve with Spark for *consistency with distributed training*. What does that cost? The *deployed forest*, exported, same board:

Engine	Flows/s	Batch p95 (ms)	mJ per flow
PySpark (deployed)	1.4	7515	1506
NumPy	54.8	194	11.3
ONNX Runtime	**10,092**	**1.2**	**0.087**

Batch 10, as deployed; identical predictions on every flow. mJ per flow: energy above the board's idle power. A Spark call costs $\approx$7 s at any batch up to 1,000, so at batch 10,000 the gap is $\approx$27$\times$ (backup).

ONNX Runtime serves the same forest $\approx$**7,200**$\times$ faster. Target for the next iteration: **train on Spark, serve with ONNX/TensorRT**.

Conclusion

What we showed

- Reproducible edge IDS: Kafka + PySpark on Jetson Orin Nano Super.
- Decoupled stages: **2.5 → 17.8 flows/s** on one board (0.32 J), **22.4** on two (0.60 J); e2e p95 ≈7 s.
- End to end: **precision 100%**, **recall 69.5%**, F1 0.82; every miss is a gate skip.
- ONNX serves the same forest ≈7,200× faster.

Limitations

- Two-node lab cluster over Wi-Fi; replayed traffic.
- No transfer to CSE-CIC-IDS2018 without retraining.
- Kafka and the database stay on a host PC.
- GPU unused (CPU-only pipeline); evasion untested.

Next

- ONNX/TensorRT serving; live traffic; retraining under drift.

Thank you — questions?

Key references

[1] Sharafaldin, I., Lashkari, A. H. & Ghorbani, A. A. *Toward generating a new intrusion detection dataset and intrusion traffic characterization*. ICISSP 2018. **CICIDS2017**

[2] Shi, W. et al. *Edge computing: Vision and challenges*. IEEE Internet of Things Journal 3(5), 2016. **edge setting**

[3] Chandola, V., Banerjee, A. & Kumar, V. *Anomaly detection: A survey*. ACM Computing Surveys 41(3), 2009. **anomaly gate**

[4] Lundberg, S. M. & Lee, S.-I. *A unified approach to interpreting model predictions*. NeurIPS 2017. **SHAP Top-30**

[5] Zaharia, M. et al. *Spark: Cluster computing with working sets*. HotCloud 2010. **Spark**

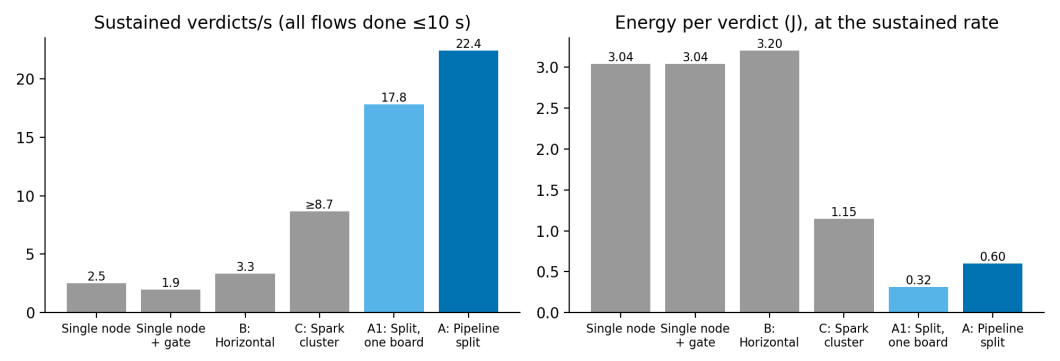
[6] Kreps, J., Narkhede, N. & Rao, J. *Kafka: a Distributed Messaging System for Log Processing*. NetDB 2011. **Kafka**

[7] Mohan Kiran Kumar, K. et al. *Distributed Intrusion Detection System using Kafka and Spark Streaming*. ICVADV 2025. **closest system**

[8] Arp, D. et al. *Dos and Don'ts of Machine Learning in Computer Security*. USENIX Security 2022. **leakage-aware protocol**

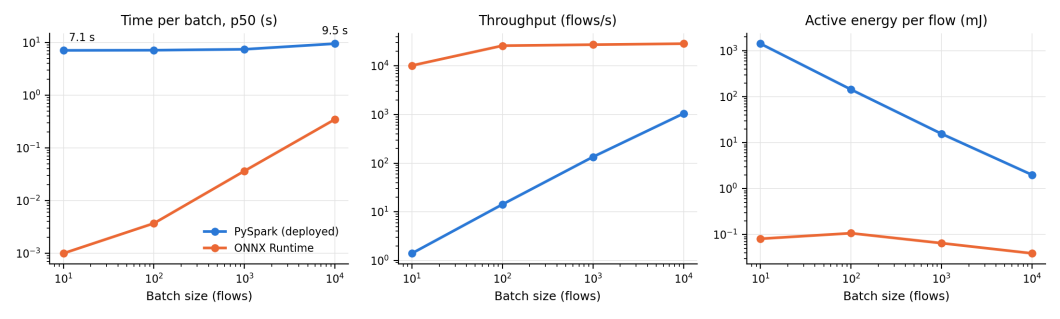
Backup slides

Backup: sustained capacity and energy per mode



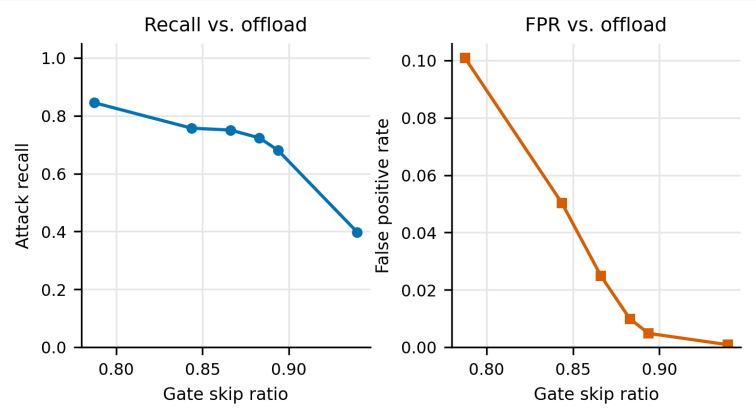
Sustained verdicts/s with every flow done within 10 s (left); energy per verdict at that rate (right).

Backup: Spark pays per call, not per tree



Same forest, same board, batch 10 to 10,000. A PySpark call takes 7.1–7.4 s up to batch 1,000 and 9.5 s at 10,000; ONNX saturates near 28,000 flows/s. The gap falls from ≈7,200× (batch 10) to ≈27× (batch 10,000).

Backup: the gate is an operating curve, not a point



Attack recall (left) and false-positive rate (right) vs. gate-skip ratio. Raising the skip ratio offloads more Spark work and costs recall; the operator chooses the point.

Backup: which model, and why

Configuration	F1	Features	Size (MB)	Edge rationale
Baseline + XGBoost	0.9971	77	2.32	Reference accuracy
SHAP Top-30 + XGBoost	0.9970	30	0.003†	Explainable; deployed feature set
RF Top-30 + XGBoost	0.9922	30	2.52	Embedded ranking
PCA k=40 + XGBoost	0.9939	40	3.82	Unsupervised reduction

We deploy **SHAP Top-30 with Random Forest** (exported model: test F1 0.9959, 30 features): a native Spark ML estimator, so the `PipelineModel` loads on ARM64 without the native libraries Spark's XGBoost integration needs. The SHAP ranking is computed after `destination_port` is removed.

†Distributed saving under-reports size; true size ≈2.5 MB.

Backup: full monitoring dashboard



Per-node throughput, detected attacks, inference latency, CPU/RAM for both boards, PostgreSQL-backed alerts.

Backup: dependence on the host PC

- Kafka, PostgreSQL, InfluxDB and Grafana run on the host: flows enter through its broker, verdicts are stored there. **Host down $\Rightarrow$ detection stops**, although both boards are healthy.
- With PySpark the host is *not* the bottleneck: the boards saturate first; host-to-board round trip 4–9 ms on Wi-Fi.
- It still bounds scaling: each added board is one more consumer of the same broker and one more writer to the same database (one insert per verdict).
- Field deployment: broker on or next to the board that receives the flow export, verdicts buffered locally and forwarded in batches; Kafka replication across the boards. **Not measured.**

Backup: measurement protocol

- Host-side orchestration starts pipeline roles on the Jetsons over SSH, drives the load, collects every verdict.
- Each verdict stores the flow's send time and replay row: latency is send to verdict, *queueing included*; skip ratio and attack recall come from the flows themselves.
- Rate sweep per mode; a rate is sustained when every flow sent in the 45 s window gets its verdict within 10 s after it. 2 runs per rate.
- Fresh Kafka consumer group per run, so no run inherits a backlog.
- Tiers batch 20 flows or flush after 1 s.
- Energy: `tegrastats`, 30 s idle baseline then load at the sustained rate; two-node modes sum both boards.